

# Import necessary library dependencies

```
In [3]: import tensorflow as tf
        from tensorflow.keras import layers, models
        from tensorflow.keras.applications import EfficientNetV2B0
        from tensorflow.keras.applications.efficientnet import preprocess_input
        from sklearn.metrics import confusion_matrix, classification_report
        import matplotlib.pyplot as plt
        import seaborn as sns
        import numpy as np
```

## Import datasets

```
In [4]: #Dataset paths
        trainpath = r'/content/drive/MyDrive/modified-dataset/train'
        validpath = r'/content/drive/MyDrive/modified-dataset/val'
        testpath = r'/content/drive/MyDrive/modified-dataset/test'
```

## Understand the Data

```
In [5]: # 1EXPLORE AND UNDERSTAND THE DATA
        IMG_SIZE = (128, 128)
        BATCH_SIZE = 32

        datatrain = tf.keras.utils.image_dataset_from_directory(trainpath, shuffle=True,
        datavalid = tf.keras.utils.image_dataset_from_directory(validpath, shuffle=True,
        datatest = tf.keras.utils.image_dataset_from_directory(testpath, shuffle=False,

        class_names = datatrain.class_names
        print(f"Classes: {class_names}")
```

Found 2410 files belonging to 10 classes.  
Found 300 files belonging to 10 classes.  
Found 310 files belonging to 10 classes.  
Classes: ['Battery', 'Keyboard', 'Microwave', 'Mobile', 'Mouse', 'PCB', 'Player', 'Printer', 'Television', 'Washing Machine']

## Visualize samples


```
In [6]: # Visualize samples
        plt.figure(figsize=(10,10))
        for images, labels in datatrain.take(1):
            for i in range(9):
                ax = plt.subplot(3,3,i+1)
```

```
plt.imshow(images[i].numpy().astype("uint8"))
plt.title(class_names[labels[i]])
plt.axis("off")
plt.show()
```

Battery



Mouse




PCB



Television



Printer



Mouse



Mouse



Mobile



Television






## Plot class distribution

```
In [7]: # Plot class distribution
def plot_class_distribution(dataset, title):
    counts = {}
    for _, labels in dataset:
        for label in labels.numpy():
            class_name = dataset.class_names[label]
            counts[class_name] = counts.get(class_name, 0) + 1
    plt.figure(figsize=(10,6))
    plt.bar(counts.keys(), counts.values(), color='skyblue')
    plt.title(title)
    plt.xticks(rotation=45)
```

```
plt.show()
```

```
plot_class_distribution(datatrain, "Training Data Distribution")  
plot_class_distribution(datavalid, "Validation Data Distribution")  
plot_class_distribution(datatest, "Test Data Distribution")
```





## DATA PREPROCESSING

```
In [8]: # DATA PREPROCESSING / PREPARATION
data_augmentation = tf.keras.Sequential([
    layers.RandomFlip("horizontal"),
    layers.RandomRotation(0.1),
    layers.RandomZoom(0.1),
])

datatrain = datatrain.prefetch(tf.data.AUTOTUNE)
datavalid = datavalid.prefetch(tf.data.AUTOTUNE)
datatest = datatest.prefetch(tf.data.AUTOTUNE)
```

## MODEL SELECTION USING EFFICIENTNET

```
In [9]: # MODEL SELECTION
base_model = EfficientNetV2B0(input_shape=IMG_SIZE+(3,), include_top=False, weights='imagenet')
for layer in base_model.layers[:100]:
    layer.trainable = False

inputs = layers.Input(shape=IMG_SIZE+(3,))
x = data_augmentation(inputs)
x = preprocess_input(x)
x = base_model(x, training=False)
```

```
x = layers.GlobalAveragePooling2D()(x)
x = layers.Dropout(0.2)(x)
outputs = layers.Dense(10, activation='softmax')(x)
model = models.Model(inputs, outputs)
```

Downloading data from [https://storage.googleapis.com/tensorflow/keras-applications/efficientnet\\_v2/efficientnetv2-b0\\_notop.h5](https://storage.googleapis.com/tensorflow/keras-applications/efficientnet_v2/efficientnetv2-b0_notop.h5)  
 24274472/24274472 ————— 0s 0us/step

## MODEL TRAINING

In [10]:

```
# MODEL TRAINING
from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau
model.compile(optimizer=tf.keras.optimizers.Adam(1e-4),
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

early_stop = EarlyStopping(monitor='val_loss', patience=3, restore_best_weights=True)
reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=2)

history = model.fit(data_train, validation_data=data_valid, epochs=15, callbacks=[early_stop, reduce_lr])
```

Epoch 1/15

76/76 ————— 51s 183ms/step - accuracy: 0.2921 - loss: 2.0460 - val\_accuracy: 0.7967 - val\_loss: 1.0489 - learning\_rate: 1.0000e-04

Epoch 2/15

76/76 ————— 10s 135ms/step - accuracy: 0.7852 - loss: 0.9561 - val\_accuracy: 0.8900 - val\_loss: 0.4801 - learning\_rate: 1.0000e-04

Epoch 3/15

76/76 ————— 21s 138ms/step - accuracy: 0.8765 - loss: 0.4865 - val\_accuracy: 0.9267 - val\_loss: 0.2935 - learning\_rate: 1.0000e-04

Epoch 4/15

76/76 ————— 20s 136ms/step - accuracy: 0.9105 - loss: 0.3234 - val\_accuracy: 0.9567 - val\_loss: 0.2241 - learning\_rate: 1.0000e-04

Epoch 5/15

76/76 ————— 20s 133ms/step - accuracy: 0.9268 - loss: 0.2303 - val\_accuracy: 0.9633 - val\_loss: 0.1834 - learning\_rate: 1.0000e-04

Epoch 6/15

76/76 ————— 9s 119ms/step - accuracy: 0.9395 - loss: 0.2093 - val\_accuracy: 0.9600 - val\_loss: 0.1685 - learning\_rate: 1.0000e-04

Epoch 7/15

76/76 ————— 11s 125ms/step - accuracy: 0.9670 - loss: 0.1370 - val\_accuracy: 0.9700 - val\_loss: 0.1556 - learning\_rate: 1.0000e-04

Epoch 8/15

76/76 ————— 10s 136ms/step - accuracy: 0.9681 - loss: 0.1139 - val\_accuracy: 0.9700 - val\_loss: 0.1572 - learning\_rate: 1.0000e-04

Epoch 9/15

76/76 ————— 20s 129ms/step - accuracy: 0.9606 - loss: 0.1249 - val\_accuracy: 0.9700 - val\_loss: 0.1444 - learning\_rate: 1.0000e-04

Epoch 10/15

76/76 ————— 9s 114ms/step - accuracy: 0.9818 - loss: 0.0712 - val\_accuracy: 0.9700 - val\_loss: 0.1392 - learning\_rate: 1.0000e-04

Epoch 11/15

76/76 ————— 10s 124ms/step - accuracy: 0.9872 - loss: 0.0675 - val\_accuracy: 0.9700 - val\_loss: 0.1392 - learning\_rate: 1.0000e-04

```

76/76 ————— 10s 134ms/step - accuracy: 0.9872 - loss: 0.0073 - val_
accuracy: 0.9667 - val_loss: 0.1409 - learning_rate: 1.0000e-04
Epoch 12/15
76/76 ————— 21s 137ms/step - accuracy: 0.9892 - loss: 0.0500 - val_
accuracy: 0.9733 - val_loss: 0.1272 - learning_rate: 1.0000e-04
Epoch 13/15
76/76 ————— 9s 121ms/step - accuracy: 0.9865 - loss: 0.0515 - val_a
ccuracy: 0.9633 - val_loss: 0.1375 - learning_rate: 1.0000e-04
Epoch 14/15
76/76 ————— 11s 145ms/step - accuracy: 0.9821 - loss: 0.0611 - val_
accuracy: 0.9600 - val_loss: 0.1519 - learning_rate: 1.0000e-04
Epoch 15/15
76/76 ————— 20s 135ms/step - accuracy: 0.9884 - loss: 0.0445 - val_
accuracy: 0.9700 - val_loss: 0.1400 - learning_rate: 5.0000e-05

```

## MODEL TUNING AND OPTIMIZATION

```

In [11]: # MODEL TUNING AND OPTIMIZATION (optional fine-tune)
for layer in base_model.layers[100:]:
    layer.trainable = True
model.compile(optimizer=tf.keras.optimizers.Adam(1e-5),
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
model.fit(datatrain, validation_data=datavalid, epochs=5, callbacks=[early_stop,

```

```

Epoch 1/5
76/76 ————— 48s 179ms/step - accuracy: 0.9864 - loss: 0.0595 - val_
accuracy: 0.9700 - val_loss: 0.1295 - learning_rate: 1.0000e-05
Epoch 2/5
76/76 ————— 18s 150ms/step - accuracy: 0.9882 - loss: 0.0549 - val_
accuracy: 0.9667 - val_loss: 0.1223 - learning_rate: 1.0000e-05
Epoch 3/5
76/76 ————— 19s 133ms/step - accuracy: 0.9836 - loss: 0.0571 - val_
accuracy: 0.9700 - val_loss: 0.1313 - learning_rate: 1.0000e-05
Epoch 4/5
76/76 ————— 11s 143ms/step - accuracy: 0.9875 - loss: 0.0465 - val_
accuracy: 0.9633 - val_loss: 0.1313 - learning_rate: 1.0000e-05
Epoch 5/5
76/76 ————— 20s 140ms/step - accuracy: 0.9919 - loss: 0.0384 - val_
accuracy: 0.9667 - val_loss: 0.1310 - learning_rate: 5.0000e-06

```

```

Out[11]: <keras.src.callbacks.history.History at 0x7e0c9268c550>

```


## MODEL PERFORMANCE VISUALIZATION

```

In [12]: # MODEL PERFORMANCE VISUALIZATION
plt.figure(figsize=(12,5))
plt.subplot(1,2,1)
plt.plot(history.history['accuracy'], label='Train Acc')
plt.plot(history.history['val_accuracy'], label='Val Acc')
plt.title('Accuracy')
plt.legend()
plt.subplot(1,2,2)
plt.plot(history.history['loss'], label='Train Loss')

```

```
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Val Loss')
plt.title('Loss')
plt.legend()
plt.show()
```



## MODEL EVALUATION AND CONFUSION MATRIX

In [13]:

```
# MODEL EVALUATION
test_loss, test_acc = model.evaluate(datatest)
print(f"Test Accuracy: {test_acc:.4f}, Test Loss: {test_loss:.4f}")

y_true = np.concatenate([y.numpy() for _, y in datatest], axis=0)
y_pred = np.argmax(model.predict(datatest), axis=1)

cm = confusion_matrix(y_true, y_pred)
print(classification_report(y_true, y_pred, target_names=class_names))

plt.figure(figsize=(10,8))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=class_names, yticklabels=class_names)
plt.xlabel("Predicted")
plt.ylabel("True")
plt.title("Confusion Matrix")
plt.show()
```


10/10 ————— 1s 135ms/step - accuracy: 0.9483 - loss: 0.1628

Test Accuracy: 0.9613, Test Loss: 0.1219

10/10 ————— 7s 462ms/step

	precision	recall	f1-score	support
Battery	0.96	0.87	0.91	30
Keyboard	1.00	1.00	1.00	30
Microwave	0.88	0.93	0.90	30
Mobile	1.00	1.00	1.00	40
Mouse	1.00	1.00	1.00	30

Mouse	1.00	1.00	1.00	30
PCB	0.94	0.97	0.95	30
Player	0.93	0.93	0.93	30
Printer	0.94	1.00	0.97	30
Television	0.96	0.90	0.93	30
Washing Machine	1.00	1.00	1.00	30
accuracy			0.96	310
macro avg	0.96	0.96	0.96	310
weighted avg	0.96	0.96	0.96	310



## FINAL TESTING AND SAVE THE MODEL

```
In [14]: # FINAL TESTING AND SAVE THE MODEL
model.save('efficientnetv2b0_ewaste_final.h5')
print("Keras model saved")
```

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save\_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my\_model.keras')` or `keras.saving.save\_model(model, 'my\_model.keras')`.

Keras model saved



# Save TFLite

```
In [15]: # Save TFLite
converter = tf.lite.TFLiteConverter.from_keras_model(model)
converter.optimizations = [tf.lite.Optimize.DEFAULT]
tflite_model = converter.convert()
with open('efficientnetv2b0_ewaste_final.tflite', 'wb') as f:
    f.write(tflite_model)
print("TFLite model saved")
```

Saved artifact at '/tmp/tmpz82u9cg5'. The following endpoints are available:

\* Endpoint 'serve'

args\_0 (POSITIONAL\_ONLY): TensorSpec(shape=(None, 128, 128, 3), dtype=tf.float32, name='keras\_tensor\_270')

Output Type:

TensorSpec(shape=(None, 10), dtype=tf.float32, name=None)

Captures:

138594976151120: TensorSpec(shape=(1, 1, 1, 3), dtype=tf.float32, name=None)  
138594976150928: TensorSpec(shape=(1, 1, 1, 3), dtype=tf.float32, name=None)  
138594976148240: TensorSpec(shape=(), dtype=tf.resource, name=None)  
138594976153232: TensorSpec(shape=(), dtype=tf.resource, name=None)  
138594976148816: TensorSpec(shape=(), dtype=tf.resource, name=None)  
138594976151504: TensorSpec(shape=(), dtype=tf.resource, name=None)  
138594976153424: TensorSpec(shape=(), dtype=tf.resource, name=None)  
138594651242960: TensorSpec(shape=(), dtype=tf.resource, name=None)  
138594651244112: TensorSpec(shape=(), dtype=tf.resource, name=None)  
138594651245072: TensorSpec(shape=(), dtype=tf.resource, name=None)  
138594651244304: TensorSpec(shape=(), dtype=tf.resource, name=None)  
138594651243536: TensorSpec(shape=(), dtype=tf.resource, name=None)  
138594976148432: TensorSpec(shape=(), dtype=tf.resource, name=None)  
138594651247184: TensorSpec(shape=(), dtype=tf.resource, name=None)  
138594651246416: TensorSpec(shape=(), dtype=tf.resource, name=None)  
138594651246608: TensorSpec(shape=(), dtype=tf.resource, name=None)  
138594651246800: TensorSpec(shape=(), dtype=tf.resource, name=None)  
138594651245648: TensorSpec(shape=(), dtype=tf.resource, name=None)  
138594651249104: TensorSpec(shape=(), dtype=tf.resource, name=None)  
138594651250064: TensorSpec(shape=(), dtype=tf.resource, name=None)  
138594651249296: TensorSpec(shape=(), dtype=tf.resource, name=None)  
138594651244496: TensorSpec(shape=(), dtype=tf.resource, name=None)  
138594651251024: TensorSpec(shape=(), dtype=tf.resource, name=None)  
138594651249488: TensorSpec(shape=(), dtype=tf.resource, name=None)  
138594651251792: TensorSpec(shape=(), dtype=tf.resource, name=None)  
138594651250256: TensorSpec(shape=(), dtype=tf.resource, name=None)  
138594651248336: TensorSpec(shape=(), dtype=tf.resource, name=None)  
138594651253328: TensorSpec(shape=(), dtype=tf.resource, name=None)  
138594651252944: TensorSpec(shape=(), dtype=tf.resource, name=None)  
138594651254288: TensorSpec(shape=(), dtype=tf.resource, name=None)  
138594651253520: TensorSpec(shape=(), dtype=tf.resource, name=None)  
138594651251408: TensorSpec(shape=(), dtype=tf.resource, name=None)  
138594651253712: TensorSpec(shape=(), dtype=tf.resource, name=None)  
138594651255632: TensorSpec(shape=(), dtype=tf.resource, name=None)  
138594651256592: TensorSpec(shape=(), dtype=tf.resource, name=None)  
138594651255824: TensorSpec(shape=(), dtype=tf.resource, name=None)  
138594651254480: TensorSpec(shape=(), dtype=tf.resource, name=None)  
138594651257552: TensorSpec(shape=(), dtype=tf.resource, name=None)  
138594651257168: TensorSpec(shape=(), dtype=tf.resource, name=None)  
138594651258704: TensorSpec(shape=(), dtype=tf.resource, name=None)

[illegible]

[illegible]

[illegible]



[illegible]

[illegible]

```

138594574603984: TensorSpec(shape=(), dtype=tf.resource, name=None)
138594574603600: TensorSpec(shape=(), dtype=tf.resource, name=None)
138594574604944: TensorSpec(shape=(), dtype=tf.resource, name=None)
138594574604176: TensorSpec(shape=(), dtype=tf.resource, name=None)
138594574601872: TensorSpec(shape=(), dtype=tf.resource, name=None)
138594574605904: TensorSpec(shape=(), dtype=tf.resource, name=None)
138594574606096: TensorSpec(shape=(), dtype=tf.resource, name=None)
138594574604368: TensorSpec(shape=(), dtype=tf.resource, name=None)
138594574606480: TensorSpec(shape=(), dtype=tf.resource, name=None)
138594574602448: TensorSpec(shape=(), dtype=tf.resource, name=None)
138594574607632: TensorSpec(shape=(), dtype=tf.resource, name=None)
138594574608400: TensorSpec(shape=(), dtype=tf.resource, name=None)
138594574607248: TensorSpec(shape=(), dtype=tf.resource, name=None)
138594574607056: TensorSpec(shape=(), dtype=tf.resource, name=None)
138594574607824: TensorSpec(shape=(), dtype=tf.resource, name=None)
138594574609744: TensorSpec(shape=(), dtype=tf.resource, name=None)
138594574610704: TensorSpec(shape=(), dtype=tf.resource, name=None)
138594574609936: TensorSpec(shape=(), dtype=tf.resource, name=None)
138594574608592: TensorSpec(shape=(), dtype=tf.resource, name=None)
138594574608976: TensorSpec(shape=(), dtype=tf.resource, name=None)
138594574613008: TensorSpec(shape=(), dtype=tf.resource, name=None)
TFLite model saved

```

## Predictions on sample test images

```

In [16]: # Show predictions on sample test images
for images, labels in datatest.take(1):
    preds = model.predict(images)
    pred_classes = tf.argmax(preds, axis=1)
    for i in range(8):
        plt.imshow(images[i].numpy().astype("uint8"))
        plt.title(f"True: {class_names[labels[i]]}, Pred: {class_names[pred_classes[i]]}")
        plt.axis("off")
        plt.show()

```

1/1 ————— 2s 2s/step

True: Battery, Pred: Battery





True: Battery, Pred: Battery



True: Battery, Pred: Battery



True: Battery, Pred: Battery





True: Battery, Pred: Battery



True: Battery, Pred: Battery





True: Battery, Pred: PCB



True: Battery, Pred: Battery



# Using CNN Model

In [17]:

```
# Normal CNN Model
normal_model = tf.keras.Sequential([
    tf.keras.layers.Rescaling(1./255, input_shape=IMG_SIZE+(3,)), # Simple rescaling
    tf.keras.layers.Conv2D(32, 3, activation='relu'),
    tf.keras.layers.MaxPooling2D(),
    tf.keras.layers.Conv2D(64, 3, activation='relu'),
    tf.keras.layers.MaxPooling2D(),
    tf.keras.layers.Conv2D(128, 3, activation='relu'),
    tf.keras.layers.GlobalAveragePooling2D(),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation='softmax')
])

normal_model.compile(optimizer='adam',
                    loss='sparse_categorical_crossentropy',
                    metrics=['accuracy'])
normal_history = normal_model.fit(data_train,
                                validation_data=data_valid,
                                epochs=15,
                                callbacks=[early_stop, reduce_lr])
```

Epoch 1/15

/usr/local/lib/python3.11/dist-packages/keras/src/layers/preprocessing/tf\_data\_layer.py:19: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(**kwargs)
```

76/76 ————— 15s 139ms/step - accuracy: 0.1358 - loss: 2.2654 - val\_accuracy: 0.1867 - val\_loss: 2.0728 - learning\_rate: 0.0010

Epoch 2/15

76/76 ————— 7s 92ms/step - accuracy: 0.2451 - loss: 2.0675 - val\_accuracy: 0.3267 - val\_loss: 2.0186 - learning\_rate: 0.0010

Epoch 3/15

76/76 ————— 10s 93ms/step - accuracy: 0.3169 - loss: 1.9402 - val\_accuracy: 0.3100 - val\_loss: 1.8574 - learning\_rate: 0.0010

Epoch 4/15

76/76 ————— 10s 96ms/step - accuracy: 0.3235 - loss: 1.8553 - val\_accuracy: 0.3733 - val\_loss: 1.7331 - learning\_rate: 0.0010

Epoch 5/15

76/76 ————— 6s 83ms/step - accuracy: 0.3490 - loss: 1.7900 - val\_accuracy: 0.3933 - val\_loss: 1.7310 - learning\_rate: 0.0010

Epoch 6/15

76/76 ————— 7s 94ms/step - accuracy: 0.3539 - loss: 1.7752 - val\_accuracy: 0.4100 - val\_loss: 1.6510 - learning\_rate: 0.0010

Epoch 7/15

76/76 ————— 10s 93ms/step - accuracy: 0.3729 - loss: 1.7356 - val\_accuracy: 0.3867 - val\_loss: 1.6837 - learning\_rate: 0.0010

Epoch 8/15

76/76 ————— 7s 94ms/step - accuracy: 0.4131 - loss: 1.6530 - val\_accuracy: 0.4433 - val\_loss: 1.6153 - learning\_rate: 0.0010  
Epoch 9/15

76/76 ————— 7s 98ms/step - accuracy: 0.4366 - loss: 1.5812 - val\_accuracy: 0.4767 - val\_loss: 1.5214 - learning\_rate: 0.0010  
Epoch 10/15

76/76 ————— 7s 94ms/step - accuracy: 0.4337 - loss: 1.5523 - val\_accuracy: 0.4600 - val\_loss: 1.4903 - learning\_rate: 0.0010  
Epoch 11/15

76/76 ————— 7s 94ms/step - accuracy: 0.4535 - loss: 1.5012 - val\_accuracy: 0.4767 - val\_loss: 1.4341 - learning\_rate: 0.0010  
Epoch 12/15

76/76 ————— 10s 87ms/step - accuracy: 0.4696 - loss: 1.4901 - val\_accuracy: 0.4967 - val\_loss: 1.4864 - learning\_rate: 0.0010  
Epoch 13/15

76/76 ————— 11s 91ms/step - accuracy: 0.4995 - loss: 1.4368 - val\_accuracy: 0.5133 - val\_loss: 1.3244 - learning\_rate: 0.0010  
Epoch 14/15

76/76 ————— 8s 101ms/step - accuracy: 0.5181 - loss: 1.3933 - val\_accuracy: 0.5600 - val\_loss: 1.2314 - learning\_rate: 0.0010  
Epoch 15/15

76/76 ————— 9s 88ms/step - accuracy: 0.5226 - loss: 1.3411 - val\_accuracy: 0.5233 - val\_loss: 1.2612 - learning\_rate: 0.0010

## Plot between CNN and EfficientNet

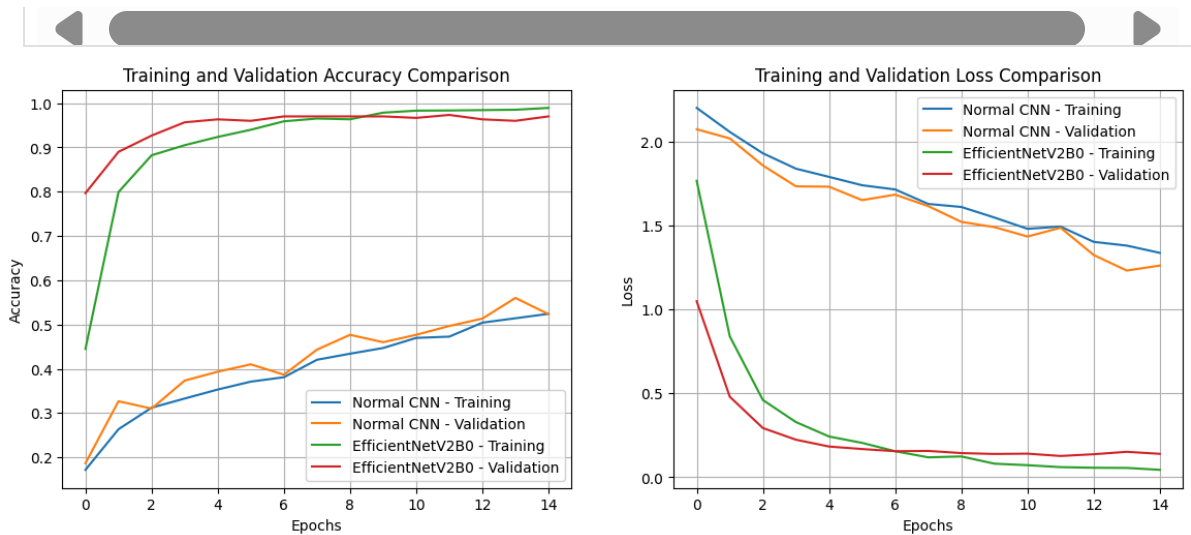
```
In [18]: import matplotlib.pyplot as plt

plt.figure(figsize=(14, 5))

# Accuracy Comparison
plt.subplot(1, 2, 1)
plt.plot(normal_history.history['accuracy'], label='Normal CNN - Training')
plt.plot(normal_history.history['val_accuracy'], label='Normal CNN - Validation')
plt.plot(history.history['accuracy'], label='EfficientNetV2B0 - Training')
plt.plot(history.history['val_accuracy'], label='EfficientNetV2B0 - Validation')
plt.title('Training and Validation Accuracy Comparison')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.grid(True)

# Loss Comparison
plt.subplot(1, 2, 2)
plt.plot(normal_history.history['loss'], label='Normal CNN - Training')
plt.plot(normal_history.history['val_loss'], label='Normal CNN - Validation')
plt.plot(history.history['loss'], label='EfficientNetV2B0 - Training')
plt.plot(history.history['val_loss'], label='EfficientNetV2B0 - Validation')
plt.title('Training and Validation Loss Comparison')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.grid(True)

plt.show()
```



```
In [20]: !pip install gradio tensorflow pillow
```

```
Requirement already satisfied: gradio in /usr/local/lib/python3.11/dist-packages (5.31.0)
Requirement already satisfied: tensorflow in /usr/local/lib/python3.11/dist-packages (2.18.0)
Requirement already satisfied: pillow in /usr/local/lib/python3.11/dist-packages (11.2.1)
Requirement already satisfied: aiofiles<25.0,>=22.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (24.1.0)
Requirement already satisfied: anyio<5.0,>=3.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (4.9.0)
Requirement already satisfied: fastapi<1.0,>=0.115.2 in /usr/local/lib/python3.11/dist-packages (from gradio) (0.115.12)
Requirement already satisfied: ffmpeg in /usr/local/lib/python3.11/dist-packages (from gradio) (0.6.0)
Requirement already satisfied: gradio-client==1.10.1 in /usr/local/lib/python3.11/dist-packages (from gradio) (1.10.1)
Requirement already satisfied: groovy~=0.1 in /usr/local/lib/python3.11/dist-packages (from gradio) (0.1.2)
Requirement already satisfied: httpx>=0.24.1 in /usr/local/lib/python3.11/dist-packages (from gradio) (0.28.1)
Requirement already satisfied: huggingface-hub>=0.28.1 in /usr/local/lib/python3.11/dist-packages (from gradio) (0.33.0)
Requirement already satisfied: jinja2<4.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (3.1.6)
Requirement already satisfied: markupsafe<4.0,>=2.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (3.0.2)
Requirement already satisfied: numpy<3.0,>=1.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (2.0.2)
Requirement already satisfied: orjson~=3.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (3.10.18)
Requirement already satisfied: packaging in /usr/local/lib/python3.11/dist-packages (from gradio) (24.2)
Requirement already satisfied: pandas<3.0,>=1.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (2.2.2)
Requirement already satisfied: pydantic<2.12,>=2.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (2.11.7)
Requirement already satisfied: pydub in /usr/local/lib/python3.11/dist-packages (from gradio) (0.25.1)
Requirement already satisfied: python-multipart>=0.0.18 in /usr/local/lib/python3.11/dist-packages (from gradio) (0.0.20)
```

Requirement already satisfied: pyyaml<7.0,>=5.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (6.0.2)

Requirement already satisfied: ruff>=0.9.3 in /usr/local/lib/python3.11/dist-packages (from gradio) (0.11.13)

Requirement already satisfied: safehttpx<0.2.0,>=0.1.6 in /usr/local/lib/python3.11/dist-packages (from gradio) (0.1.6)

Requirement already satisfied: semantic-version~=2.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (2.10.0)

Requirement already satisfied: starlette<1.0,>=0.40.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (0.46.2)

Requirement already satisfied: tomlkit<0.14.0,>=0.12.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (0.13.3)

Requirement already satisfied: typer<1.0,>=0.12 in /usr/local/lib/python3.11/dist-packages (from gradio) (0.16.0)

Requirement already satisfied: typing-extensions~=4.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (4.14.0)

Requirement already satisfied: uvicorn>=0.14.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (0.34.3)

Requirement already satisfied: fsspec in /usr/local/lib/python3.11/dist-packages (from gradio-client==1.10.1->gradio) (2025.3.2)

Requirement already satisfied: websockets<16.0,>=10.0 in /usr/local/lib/python3.11/dist-packages (from gradio-client==1.10.1->gradio) (15.0.1)

Requirement already satisfied: absl-py>=1.0.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (1.4.0)

Requirement already satisfied: astunparse>=1.6.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (1.6.3)

Requirement already satisfied: flatbuffers>=24.3.25 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (25.2.10)

Requirement already satisfied: gast!=0.5.0,!0.5.1,!0.5.2,>=0.2.1 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (0.6.0)

Requirement already satisfied: google-pasta>=0.1.1 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (0.2.0)

Requirement already satisfied: libclang>=13.0.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (18.1.1)

Requirement already satisfied: opt-einsum>=2.3.2 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (3.4.0)

Requirement already satisfied: protobuf!=4.21.0,!4.21.1,!4.21.2,!4.21.3,!4.21.4,!4.21.5,<6.0.0dev,>=3.20.3 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (5.29.5)

Requirement already satisfied: requests<3,>=2.21.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (2.32.3)

Requirement already satisfied: setuptools in /usr/local/lib/python3.11/dist-packages (from tensorflow) (75.2.0)

Requirement already satisfied: six>=1.12.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (1.17.0)

Requirement already satisfied: termcolor>=1.1.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (3.1.0)

Requirement already satisfied: wrapt>=1.11.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (1.17.2)

Requirement already satisfied: grpcio<2.0,>=1.24.3 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (1.73.0)

Requirement already satisfied: tensorboard<2.19,>=2.18 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (2.18.0)

Requirement already satisfied: keras>=3.5.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (3.8.0)

Requirement already satisfied: h5py>=3.11.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (3.14.0)

Requirement already satisfied: ml-dtypes<0.5.0,>=0.4.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (0.4.1)

Requirement already satisfied: tensorflow-io-gcs-filesystem>=0.23.1 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (0.37.1)



```

110/ python3.11/dist-packages (from tensorflow==2.15.0)
Requirement already satisfied: idna>=2.8 in /usr/local/lib/python3.11/dist-packages (from anyio<5.0,>=3.0->gradio) (3.10)
Requirement already satisfied: sniffio>=1.1 in /usr/local/lib/python3.11/dist-packages (from anyio<5.0,>=3.0->gradio) (1.3.1)
Requirement already satisfied: wheel<1.0,>=0.23.0 in /usr/local/lib/python3.11/dist-packages (from astunparse>=1.6.0->tensorflow) (0.45.1)
Requirement already satisfied: certifi in /usr/local/lib/python3.11/dist-packages (from httpx>=0.24.1->gradio) (2025.6.15)
Requirement already satisfied: httpcore==1.* in /usr/local/lib/python3.11/dist-packages (from httpx>=0.24.1->gradio) (1.0.9)
Requirement already satisfied: h11>=0.16 in /usr/local/lib/python3.11/dist-packages (from httpcore==1.*->httpx>=0.24.1->gradio) (0.16.0)
Requirement already satisfied: filelock in /usr/local/lib/python3.11/dist-packages (from huggingface-hub>=0.28.1->gradio) (3.18.0)
Requirement already satisfied: tqdm>=4.42.1 in /usr/local/lib/python3.11/dist-packages (from huggingface-hub>=0.28.1->gradio) (4.67.1)
Requirement already satisfied: hf-xet<2.0.0,>=1.1.2 in /usr/local/lib/python3.11/dist-packages (from huggingface-hub>=0.28.1->gradio) (1.1.3)
Requirement already satisfied: rich in /usr/local/lib/python3.11/dist-packages (from keras>=3.5.0->tensorflow) (13.9.4)
Requirement already satisfied: namex in /usr/local/lib/python3.11/dist-packages (from keras>=3.5.0->tensorflow) (0.1.0)
Requirement already satisfied: optree in /usr/local/lib/python3.11/dist-packages (from keras>=3.5.0->tensorflow) (0.16.0)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.11/dist-packages (from pandas<3.0,>=1.0->gradio) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.11/dist-packages (from pandas<3.0,>=1.0->gradio) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-packages (from pandas<3.0,>=1.0->gradio) (2025.2)
Requirement already satisfied: annotated-types>=0.6.0 in /usr/local/lib/python3.11/dist-packages (from pydantic<2.12,>=2.0->gradio) (0.7.0)
Requirement already satisfied: pydantic-core==2.33.2 in /usr/local/lib/python3.11/dist-packages (from pydantic<2.12,>=2.0->gradio) (2.33.2)
Requirement already satisfied: typing-inspection>=0.4.0 in /usr/local/lib/python3.11/dist-packages (from pydantic<2.12,>=2.0->gradio) (0.4.1)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.21.0->tensorflow) (3.4.2)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.21.0->tensorflow) (2.4.0)
Requirement already satisfied: markdown>=2.6.8 in /usr/local/lib/python3.11/dist-packages (from tensorboard<2.19,>=2.18->tensorflow) (3.8)
Requirement already satisfied: tensorboard-data-server<0.8.0,>=0.7.0 in /usr/local/lib/python3.11/dist-packages (from tensorboard<2.19,>=2.18->tensorflow) (0.7.2)
Requirement already satisfied: werkzeug>=1.0.1 in /usr/local/lib/python3.11/dist-packages (from tensorboard<2.19,>=2.18->tensorflow) (3.1.3)
Requirement already satisfied: click>=8.0.0 in /usr/local/lib/python3.11/dist-packages (from typer<1.0,>=0.12->gradio) (8.2.1)
Requirement already satisfied: shellingham>=1.3.0 in /usr/local/lib/python3.11/dist-packages (from typer<1.0,>=0.12->gradio) (1.5.4)
Requirement already satisfied: markdown-it-py>=2.2.0 in /usr/local/lib/python3.11/dist-packages (from rich->keras>=3.5.0->tensorflow) (3.0.0)
Requirement already satisfied: pygments<3.0.0,>=2.13.0 in /usr/local/lib/python3.11/dist-packages (from rich->keras>=3.5.0->tensorflow) (2.19.1)
Requirement already satisfied: mdurl~0.1 in /usr/local/lib/python3.11/dist-packages (from markdown-it-py>=2.2.0->rich->keras>=3.5.0->tensorflow) (0.1.2)

```

```

In [22]: import gradio as gr
import tensorflow as tf
import numpy as np

```

```

import numpy as np
from PIL import Image
from tensorflow.keras.applications.efficientnet import preprocess_input

# Load your saved Keras model
model = tf.keras.models.load_model('efficientnetv2b0_ewaste_final.h5')
class_names = [
    'Battery', 'Keyboard', 'Microwave', 'Mobile',
    'Mouse', 'PCB', 'Player', 'Printer',
    'Television', 'Washing Machine'
]

# Prediction function
def classify_image(img):
    """Preprocess input image and make prediction."""
    img = img.resize((128, 128))
    img_array = np.asarray(img, dtype=np.float32)
    img_array = preprocess_input(img_array)
    img_array = np.expand_dims(img_array, axis=0)

    prediction = model.predict(img_array)
    index = np.argmax(prediction)
    class_name = class_names[index]
    confidence = prediction[0][index]

    return f"Predicted: {class_name} (Confidence: {confidence:.2f})"

# Create Gradio Interface
iface = gr.Interface(
    fn=classify_image,
    inputs=gr.Image(type="pil"),
    outputs="text",
    title="E-Waste Classifier",
    description="Upload an image of an e-waste item for classification."
)

# Launch the app
iface.launch(share=True) # Enables a public link

```

WARNING:absl:Compiled the loaded model, but the compiled metrics have yet to be built. `model.compile\_metrics` will be empty until you train or evaluate the model. Colab notebook detected. To show errors in colab notebook, set debug=True in launch().

\* Running on public URL: <https://80de40e225bc483c3d.gradio.live>

This share link expires in 1 week. For free permanent hosting and GPU upgrades, run `gradio deploy` from the terminal in the working directory to deploy to Hugging Face Spaces (<https://huggingface.co/spaces>)

Out[22]: